Five specialist roles. One optional local LLM.

A transparent AI workflow that explains business evidence without giving a model database permissions.

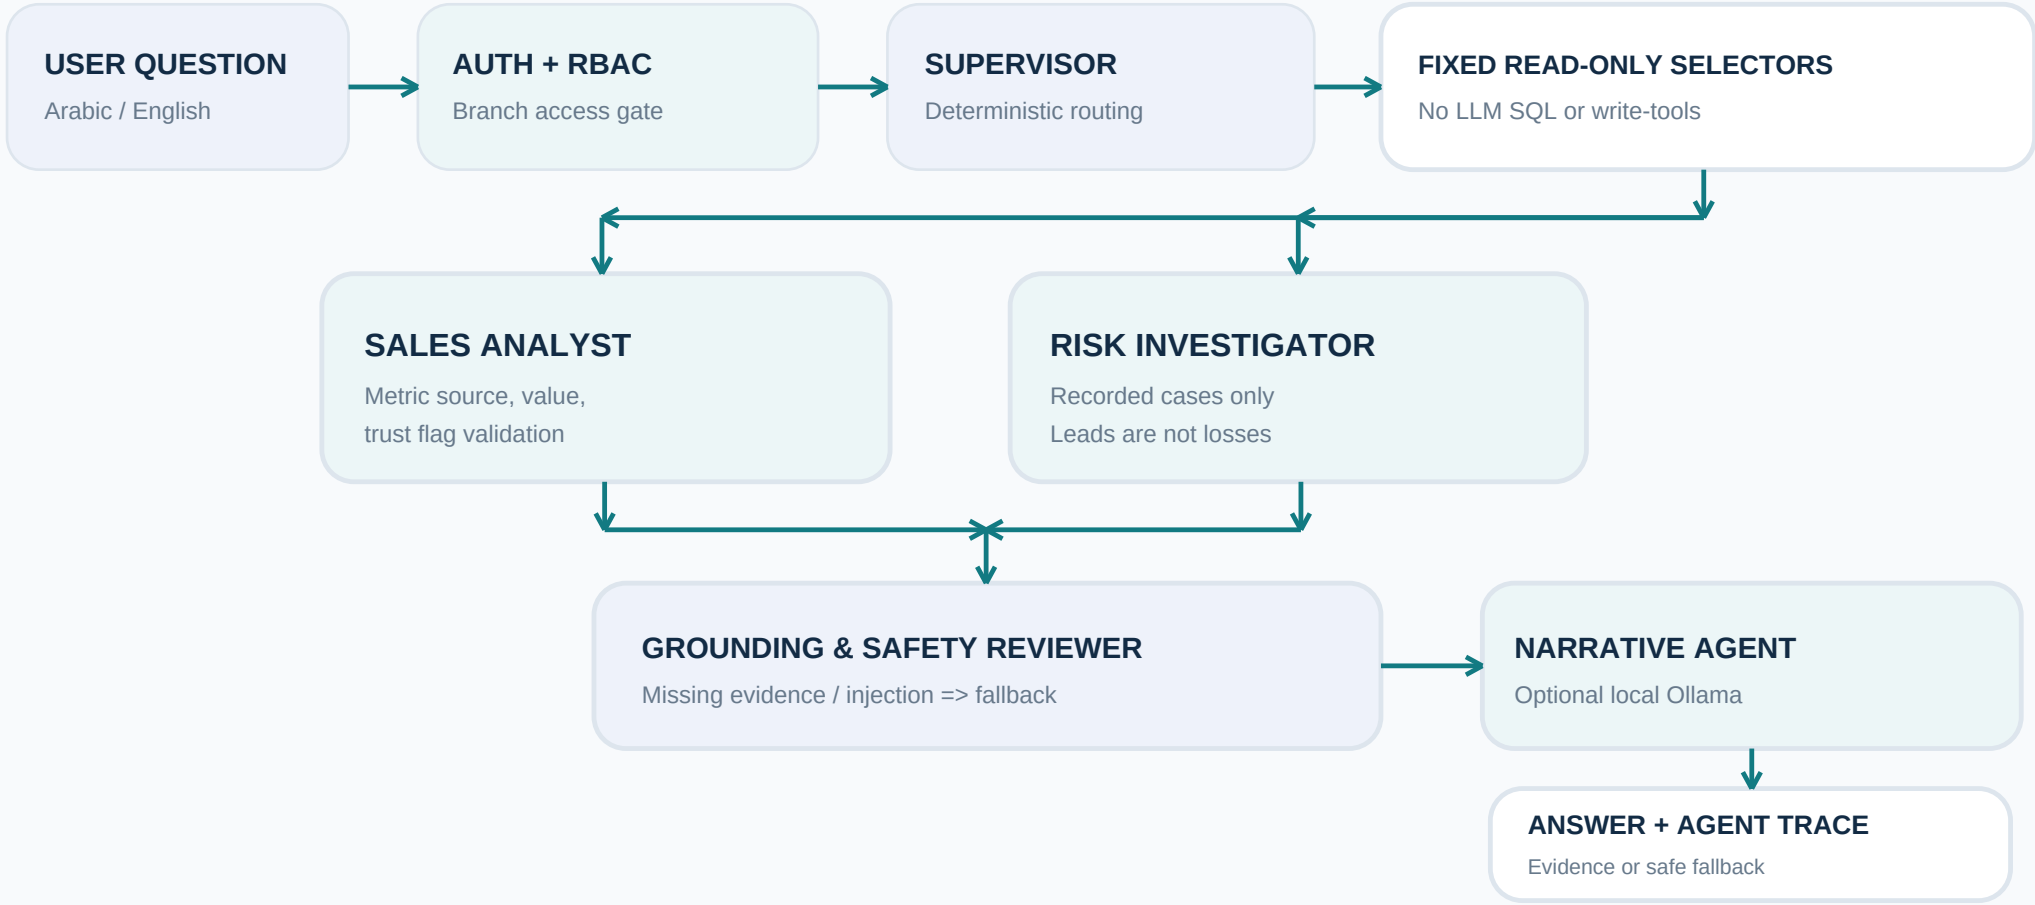

Who does what - and what each role cannot do

The implementation uses role-specific Python functions. It is not an autonomous five-model swarm.

01	**SUPERVISOR**	Selects bounded intent and data type	**NEVER** Cannot grant privileges or run SQL
02	**SALES ANALYST**	Validates value, source and trust	**NEVER** Cannot invent company metrics
03	**RISK INVESTIGATOR**	Reviews actual recorded case evidence	**NEVER** Cannot assert unproven losses
04	**GROUNDING REVIEWER**	Decides when LLM use is permitted	**NEVER** Cannot bypass authentication
05	**NARRATIVE AGENT**	Produces bilingual grounded explanation	**NEVER** Cannot call write tools or issue SQL

The backend owns authorization and evidence. Ollama owns wording only, and is optional.

Only fixed read-only evidence paths are available; no dynamic agent tools or database write permissions.

Designed to be demonstrated - and checked

1. LIVE REQUEST CONTRACT

POST /api/copilot/chat

Response includes answer + sources + agent_trace

mode: grounded-llm | deterministic-fallback

llm_used: true ONLY when generation succeeded

GET /api/copilot/agents - read-only metadata

2. VERIFIED VS. NOT VERIFIED

PASS - 113 / 113 backend tests

PASS - changed TSX syntax-level check

BLOCKED - frontend install / build / browser E2E

NOT RUN - live Ollama / Windows runtime

NO-GO - enterprise production readiness

3. REPRODUCIBLE PORTFOLIO WALKTHROUGH

01	Generate separate synthetic demo DB	py scripts\seed_synthetic_demo.py
02	Enable local Ollama	scripts\enable_ollama_demo.ps1
03	Ask about synthetic net sales	Sales Analyst -> Reviewer -> Narrative
04	Ask about recorded investigations	Risk Investigator -> Reviewer -> Narrative
05	Ask for an unsupported exact financial loss	Reviewer blocks fabrication -> fallback

Proof matters: show the actual expandable workflow, data provenance, and honest model status.